

— KOICA ODA · NIGERIA · ABUJA

Nigeria Startup Digital Innovation

Expectations in the AI Era.

Two a16z primary charts · Lean Leaderboard · Big-tech restructuring · The unemployment flip

5th Dispatch · A 30-minute briefing — then the real class begins.

\$8.9M

Consumer Top 25% · Year-1
ARR

\$5.3M

Enterprise Top 25% · Year-1
ARR

36%

Solo-founded share 2025 · vs
17% (2016)

5.8%

U.S. recent-grad unemployment
· flips national 4.2%

INSTRUCTOR

KyungBae Park (Ph.D.) — CEO, Anvyl Labs

Adjunct Professor, Korea University Business School

LOCATION

Abuja, Nigeria

Why solo and micro-team AI startups suddenly own the **U.S. VC conversation**


"There will soon be a one-person, billion-dollar company."
— Sam Altman, OpenAI

That remark set off **leanaileaderboard.com**, built by **Henry Shi** (co-founder, Super.com) and since covered by WSJ, NYT, Forbes and The Information. The primary data behind it: over nine years the **solo-founder share of new U.S. startups rose 17% → 36%, a 2.1× increase** (Carta, 54,501 startups).


SPEEDRUN NEWSLETTER

Share of businesses crossing \$1M within a year of going live on Stripe — the 2025 cohort ran about **30% higher** than 2023 and roughly **3× higher** than 2019.

Sources: [a16z](#) [LinkedIn](#) · [Carta](#) · [leanaileaderboard.com](#) · [Speedrun](#)



Solo-Founded Startups Are Becoming More Common




Consumer Gen AI · Year-1 ARR Ramp

Median **\$4.1M** by Month 12 / top quartile **\$8.9M** — even the median is **4.1×** the old dream benchmark of \$1M. The top quartile closes Series A in **5 months** and raises only \$3.1M before it.

Year 1 quartile	Month 6	Month 12	Time to Series A	Pre-A Raised
Bottom quartile	\$1.3M	\$2.9M	12 mo	\$10.5M
Median	\$2.5M	\$4.1M	8 mo	\$8.0M
Top quartile	\$3.8M	\$8.9M	5 mo	\$3.1M

Chart as published: Source: A16z, internal data, June 2025 · → **Speed is the primary indicator of capital efficiency**. The top quartile raises the least and reaches Series A the fastest.

ARR Ramp of Consumer Gen AI startup in Year 1



Enterprise Gen AI · Year-1 ARR Ramp

B2B startups also reach a median **\$2.1M** in year one — **2.1× the old "best-in-class \$1M"** bar. The top quartile hits \$5.3M and Series A in 7 months. → **"\$1M ARR in Year 1" is no longer remarkable.**


Year 1 quartile	Month 6	Month 12	Time to Series A	Pre-A Raised
Bottom quartile	\$0.5M	\$1.2M	13 mo	\$5.5M
Median	\$0.7M	\$2.1M	9 mo	\$4.0M
Top quartile	\$2.0M	\$5.3M	7 mo	\$2.3M

"Pre-AI, a common benchmark for best-in-class enterprise startups was \$1 million in ARR in its first 12 months. The median enterprise AI startup now hits \$2.1M ARR by Month 12 — **more than double** the old benchmark." — a16z, verbatim

ARR Ramp of Enterprise Gen AI startups in Year 1

ARR Ramp of Enterprise Gen AI startups in Year 1

Source: A16z, internal data, June 2025



The Path to Series A

	Month 6 Revenue	Month 12 Revenue	Time to Series A	Pre-Series A \$ Raised
Bottom quartile	\$0.5M	\$1.2M	13 Months	\$5.5M
Median	\$0.7M	\$2.1M	9 Months	\$4.0M
Top quartile	\$2.0M	\$5.3M	7 Months	\$2.3M

Recent-graduate unemployment flips the national average — **a first since 1990**

RECENT GRADUATES

5.8%

NATIONAL AVERAGE


4.2%

Since 1990 college graduates have *always* had lower unemployment than the national average. **From 2022 onward, graduates have it harder than the average worker.** The parent generation still says "get a degree and a good job follows" — the data says otherwise.

Teaching point: the message shifts from "study more" to "the ability to create value, and to execute."

Recent graduates vs National unemployment, 1990–2024

Experienced college graduate unemployment rates have always been lower than the national average, however not from 2022



Human Resource → Hybrid Resource

Shopify's "100× output" declaration

*"Prove AI cannot do the job
before you ask for headcount."*

— Tobi Lütke, CEO of Shopify · internal memo (Apr 2025)

ORIGINAL TWEET

"Reflexive AI usage is now a baseline expectation."

THE NEW NORMAL

Practitioners already report **100× output** with AI. Across SaaS, "**No Net New Headcount**" has become the new normal.

From Human Resource to Hybrid Resource

From Human Resource To Hybrid Resource

CEO at Shopify, Tobi Lütke said employees should demonstrate that their jobs cannot be done by AI before requesting additional headcount. He has also mentioned seeing individuals accomplish tasks previously deemed impossible by leveraging AI, resulting in a 100x increase in work output

- Not everyone can get to \$50m ARR with 5 employees
- But at scale, “No Net New Headcount” is the new normal
 - Do it first with AI
 - Do it second with an upgrade
 - Hire AI services



What This Means

1. **Using AI effectively is now a fundamental expectation of everyone at Shopify.** It's a tool of all trades today, and will only grow in importance. Frankly, I don't think it's feasible to opt out of learning the skill of applying AI in your craft; you are welcome to try, but I want to be honest I cannot see this working out today, and definitely not tomorrow. Stagnation is almost certain, and stagnation is slow-motion failure. If you're not climbing, you're sliding.
2. **AI must be part of your GSD Prototype phase.** The prototype phase of any GSD project should be dominated by AI exploration. Prototypes are meant for learning and creating information. AI dramatically accelerates this process. You can learn to produce something that other team mates can look at, use, and reason about in a fraction of the time it used to take.

4. **Learning is self directed, but share what you learned.** You have access to as much of the cutting edge AI tools as possible. There is [chat.shopify.io](#), which we had for years now. Developers have [proxy](#), Copilot, Cursor, Claude code, all pre-tocled and ready to go. We'll learn and adapt together as a team. We'll be sharing Ws (and Ls!) with each other as we experiment with new AI capabilities, and we'll dedicate time to AI integration in our monthly business reviews and product development cycles. Slack and Vault have lots of places where people share prompts that they developed, like [#revenue-ai-use-cases](#) and [#ai-centaurs](#).
5. **Before asking for more Headcount and resources,** teams must demonstrate why they cannot get what they want done using AI. What would this area look like if autonomous AI agents were already part of the team? This question can lead to really fun discussions and projects.

Agentic AI has begun to **outperform humans** in coding and in sales

MICROSOFT

30%

of code written by AI (Nadella)
9,000 laid off the same month

GOOGLE

Jules

Coding agent released **free**
The agent race begins in earnest

SAASTR · JASON LEMKIN

2024.04 "No way."

2025.04 "Maybe."

2026.04 "Better than 80% of sales reps."

Pulls expert automation **2 years forward**

Coding and sales — a startup's two core functions — are being agentified at once. "Hire more people" is dropping out of the default playbook.


Agentic AI is outperforming on software coding & sales

Agentic AI is outperforming on software codings & sales

Satya Nadella says as much as 30% of Microsoft code is written by AI
Microsoft laying off about 9,000 employees in latest round of cuts

PUBLISHED WED, JUL 2 2025·9:07 AM EDT | UPDATED 6 HOURS AGO

**GOOGLE JULES: THE ULTIMATE
 FREE AI CODING AGENT THAT'S
 CHANGING SOFTWARE
 DEVELOPMENT**



Jason
 @jasonlk

Saastr 2025 is May 13-15 Lemkin

April 2024: AI can't beat a good sales rep. No way.

April 2025: AI might beat an OK sales rep. Maybe.

April 2026: Our AI is better than 80% of sales reps

10:57 AM · Apr 29, 2025

AI layoffs begin — startup hiring has stopped too

2025 BIG-TECH RESTRUCTURING

Microsoft	6,000
Dell	12,000
IBM HR	Hundreds
PwC	1,500
Salesforce	~1,000
Klarna	~700
Chegg · Duolingo	22% · 10%

THE MORE DECISIVE SIGNAL · CARTA "STARTUP HEADCOUNT IS FROZEN"

2019–2022 · 48 months
+865K

→

2023–2025 · 27 months
+31K

Net hiring growth is down **96%** · **12 of 27 months (44%)** were net negative. Based on **3.8M** job changes within Carta companies (Q1 2019 – Q1 2025).

AI layoffs begin — Startup headcount is frozen

AI layoffs begin.....


Tech layoffs in 2025



Startup headcount is frozen

Net headcount (hires minus departures) struggling to grow in recent years

Data: Aggregate of 3.8 million job changes within Carta companies from Q1 2019—Q1 2025



Subscribe for more startup data: carta.com/data

Tiny teams are the future.

Speed is the new moat.

Company	Revenue	Time taken	Team
Cursor	\$100M ARR	21 months	20
Bolt	\$20M ARR	2 months	15
Lovable	\$10M ARR	2 months	15
Mercor	\$50M	2 years	30
ElevenLabs	\$100M ARR	2 years	50
Genspark	\$36M ARR	45 days	lean + AI agents

PEOPLE NEEDED FOR \$100M ARR

15–50

is all it takes

Caveat: existing B2B customers stay put, but **new customers expect far more** — that is precisely what "speed is the new moat" means.

The same AI creates **opportunity** and **risk** at once

OPPORTUNITY

The bar itself has moved — a16z Year-1 median ARR of \$2.1M (enterprise). A single founder can now run global SaaS.

Speed is the moat — Cursor in 21 months · Bolt \$20M in 2 months · Genspark \$36M in 45 days

Speedrun-style funding exists — a16z wires up to \$1M over 12 weeks, plus \$10M+ in credits

AI literacy + execution is the new credential — *stronger protection than a degree*

RISK · ETHICAL BALANCE

The degree shields less — U.S. recent-grad unemployment 5.8% vs national 4.2% (since 2022)

Simultaneous restructuring — Microsoft 6,000 · Dell 12,000 · IBM HR at scale (2025)

Startup hiring is frozen — Carta +865K → +31K; 12 of 27 months net negative

VC theses have converged — non-AI narratives are effectively shut out of funding

Consolidated Source Index — a16z · solo founders · Lean

#	Source	Used for
1	a16z — Revenue benchmarks for AI apps	Consumer + Enterprise ARR ramp charts ★
2	a16z — Top 100 Gen AI Consumer Apps, 6th Ed.	Consumer supplementary data
3	a16z LinkedIn — Solo-founded startups are becoming more common	17% → 36% chart
4	a16z Speedrun — The case for (and against) founding solo	Solo founding · Stripe 30% data
5	speedrun.a16z.com/faq	30-minute VC meeting FAQ
6	leanaileaderboard.com	Henry Shi's official leaderboard
7	GitHub — official-lean-ai-native-leaderboard	Leaderboard dataset
8	carta.com/data/solo-founders-report/	Nine-year 17% → 36% raw data
9	Forbes — The Race To Create A Billion-Dollar One-Person Business	Secondary coverage
10	SaaStr — Median enterprise AI startup hits \$2.1M ARR by month 12	Enterprise secondary verification
11	HBS RIS Working Paper 26-090	Academic citation of a16z ARR
12	Gennaro Cuofano (LinkedIn) — Internal a16z data	Enterprise secondary verification

Consolidated Source Index — unemployment · Shopify · Agentic · Layoffs

#	Source	Used for
13	Oxford Economics — Educated but unemployed (PDF)	Original report on the unemployment flip
14	Oxford Economics — resource page	Report landing page
15	BLS — Unemployment rates by educational attainment	Official supplement
16	NY Fed — Labor Market for Recent College Graduates	Background on 5.8% vs 4.2%
17	Encoura — Labor Market Part 1: No More Jobs	Secondary visualization
18	EPI — Class of 2026 faces a weaker labor market	Secondary visualization
19	Randal Olson — Recent-grad unemployment flip (Jun 2026)	Secondary visualization
20	X — @tobi: "Reflexive AI usage is now a baseline expectation"	Original Shopify memo
21	CNBC (2025.4.7) — Prove AI can't do jobs before asking for headcount	Shopify primary coverage
22	Forrester — analysis of the Shopify CEO memo	Shopify commentary
23	CNBC (2025.4.29) — Nadella: 30% of Microsoft code written by AI	30% of Microsoft code
24	Business Insider — Microsoft cuts 9,000 jobs (Jul 2025)	Microsoft layoffs
25	jules.google	Google Jules official
26	carta.com/data — Startup headcount is frozen	3.8M job-change dataset

02

Part 2 Section 2 of 3 · 14 slides

AI-Era One-Person Startup.

Empirical case study.

The numbers you just saw, made by three founders you've never heard of.
Then: the cost curve, an honest word, and our two tools.

Cold open

3 hero cases

Cost 99% ↓

Honesty

Codex & Antigravity

One folder

● 5-MINUTE CLIP

First, watch this.

Sam Altman on GPT-5 and the 1-person company

youtube.com/watch?v=hmtuvNfytjM&t=1280s

Starts at **21:20** — “It’s 2035. What new jobs exist?” · play ≈ 5 min (to ≈ 26:00)

*“There will soon be a
one-person, billion-dollar
company.”*

— Sam Altman, OpenAI

For the first time in history, **this is not a joke.**

You are the **luckiest generation** to build a company.

The tools that once needed 100 people now fit in your laptop.

Let's meet the people already doing it. →

● LIVE · RIGHT NOW

leanaileaderboard.com

OFFICIAL LEAN AI NATIVE COMPANIES

01	Medvi	GLP-1 telehealth	\$401M ARR
02	Pieter Levels	multi-product	\$3M+ ARR
03	Higgsfield	AI video	\$300M ARR
04	... 47 more companies		

RUN BY

Henry Shi

Super.com founder · \$150M+ exit

INCLUSION CRITERIA

- ✓ \$5M+ ARR
- ✓ Under 50 employees
- ✓ Under 5 years old

SOURCE

leanaileaderboard.com — as of 2026.04

1 MEDVI

\$20K became \$1.8B
in two years.

FOUNDER

Matthew Gallagher

Age 41

EMPLOYEES

1 / 2

⚠ sources conflict

INITIAL CAPITAL

\$20,000

bootstrapped · no VC

2025 REVENUE

\$401M

20,000× the initial capital

2026 RUN-RATE TARGET

\$1.8B

daily \$3M+ · 16.2% net margin

WHAT IT IS

GLP-1 weight-loss telehealth.

HOW

12+ AI tools
Built in 2 years.

NET INCOME

~\$65M / year

2 PIETER LEVELS

0 employees. 0 funding.
Multi-product. **\$3M+ ARR.**

PROFILE

Dutch · self-taught dev
Briefly hired one AI dev for Photo AI.

FUNDING

\$0
100% bootstrapped

PHOTO AI MARKETING SPEND

\$0

PORTFOLIO

Nomad List	community	\$700K ARR
Remote OK	job board	\$2M+ ARR
Photo AI	\$132K MRR	\$1.58M ARR
Interior AI	\$40K MRR	\$480K ARR

70 failures. 1 Photo AI **hit**.

● 65 misses ● 5 hits at \$100K+ ARR

He **publicly tracks** his failures.
Of 70+ projects, only 4–5 reached \$100K+ ARR.

"Solo maker rule = coding skill ≠ try rate × distribution."

Try rate is the lever. Multi-product creates compounding.

3 HIGGSFIELD

232 employees. **\$300M ARR.**
A different definition of *lean*.

FOUNDER

Alex Mashrabov

ex-Snap · Generative AI
Director

FOUNDED

2023

Yerzat Dulat · Mahi de Silva

EMPLOYEES

232 ⚠️

Does NOT meet
leanaileaderboard ≤50

FUNDING ·
VALUATION

\$138M · \$1.3B

2026.01

REVENUE PER EMPLOYEE


\$1.29M

Traditional SaaS average: ~\$200K

2026 EOY ARR TARGET

\$1B

2015 → 2026. What changed?



One-person startup initial cost · 2015 → 2026

~~\$70,000~~

→

\$700

3-month cumulative · same scope

Area	2015	2026
UI / UX design	\$400 — \$15K	Free — \$30/mo
Full-stack dev	\$20K+ outsourced	\$40 — \$200/mo
Server / infra	\$400 — \$1.5K/mo	Free — \$20/mo
Video / landing	\$3.5K — \$7K	\$30 — \$200/mo
Customer support	\$1.8K — \$2.5K/mo	\$7 — \$70/mo
Legal drafts	\$400 — \$1.5K/case	\$0 — \$200/mo
3-month total	\$70K — \$220K	\$200 — \$1K

≈ 99% reduction.

A full-stack company for 3 months on the price of a single laptop.

ACE AI Startup Bootcamp — Eduniety Lab

Then why isn't *everyone* doing it?

Item	2015	2026	Change
Learning curve	1+ year	1 – 3 months	↓ decreased
Cost of failure	\$70K lost	\$700 lost	↓ decreased
CAC customer acquisition	similar	higher	↑ increased
Differentiation	"existence = differentiation"	"buried in noise"	↑ increased

*Build cost converges to zero. **Sell cost stays.** Differentiation cost increased.*

Many great AI tools exist.

We use **two**.

THIS BOOTCAMP

AI CODING AGENT

Codex

Desktop app — plain English, no commands

AGENTIC IDE

Antigravity

Full-stack workspace · build in the editor

You **won't** type commands.
You **will** still read code.

OTHER GREAT OPTIONS


Tool	Sweet spot
Codex OpenAI	Desktop coding agent — our pick
Antigravity Google	Agentic full-stack IDE — our pick
Claude Code	Planning · logic · docs
Cursor	VS Code + AI
v0 Vercel	UI prototyping
Bolt.new	Full-stack in 1 minute
Replit	Instant deploy

THE SETUP

Two agents. One **folder**.

Codex and Antigravity both read and write the **same project folder** on your laptop. The folder is the single source of truth — the agents are interchangeable hands.

- 01 **Build** in Codex → writes `spec.md` and code into the folder
- 02 **Run & fix** in Antigravity → same folder, nothing copied



03

Part 3

Section 3 of 3 • 33 slides

GStack Workshop.

Forge before you build.

Install Codex or Antigravity, then think like a CEO, an engineer and a QA — with an AI team. Garry Tan's 28 slash commands, hands-on.

Choose your path

Install

What is GStack

Office Hours

CEO & Eng review

Missions 01–04

Garbage in, garbage out

Start with Why

01

Introduction Section 1 of 8

Getting Started

Two paths. One forge. Choose your tool, install it, and meet GStack.

Choose Your Path

Select the foundation for your AI engineering team. Both paths lead to GStack mastery.

PATH A

Codex

- Optimized for OpenAI models & ecosystem
- Desktop app for Mac & Windows
- Built-in multi-agent sandboxing

PATH B

Antigravity

- Google's standalone AI engineering platform
- Deep MCP (Model Context Protocol) integration
- Advanced worktree and artifact management

PREREQUISITES

Team account credentials for Codex / Antigravity — provided by the organizers.

Mac or Windows laptop with internet access.

SUPPORT

7 AI ToT Tutors available for both paths.

Raise your hand anytime during setup!

🕒 2 hours

📊 30 slides

👥 7 Tutors

github.com/garrytan/gstack • ★ 70,000+

02

Path A Section 2 of 8

Install Codex

Prepare the anvil. Deploy the OpenAI command center.

Download Codex Desktop App

Get the hammer: openai.com/codex

macOS

- 01 Download the .dmg file
- 02 Optimized for Apple Silicon
- 03 Drag to Applications folder

🗦 Windows

- 01 Download the .exe installer
- 02 Windows 10/11 supported
- 03 Follow the setup wizard



AI Tutors on Standby — Need help choosing your version? Raise your hand and your designated AI ToT will assist you immediately.

Launch & Verify

Ignite the forge. Ensure your command center is operational.

01

Open Codex Desktop App

Locate the app in your Applications or Start menu and launch it.

02

Sign in with OpenAI Account

Use the team account credentials provided by the organizers to authenticate the local environment.

03

Verify Installation

Check the bottom-left corner of the dashboard for the green **"Ready"** status indicator.



AI Tutors on Standby — Authentication failed? Status offline? Raise your hand and an AI ToT will assist you.

03

Path B

Section 3 of 8

Install Antigravity

Google's forge for rapid prototyping.

03 · PATH B

Download Google Antigravity

Get the hammer: antigravity.google/download

macOS

- 01 Apple Silicon & Intel
- 02 macOS v12+ required
- 03 Drag to Applications

🖥 Windows

- 01 Download .exe installer
- 02 Windows 10/11 64-bit
- 03 Follow setup wizard

🐧 Linux

- 01 Debian/Ubuntu supported
- 02 glibc ≥ 2.28 required
- 03 Standalone application



AI Tutors on Standby — Need help with your specific OS? Raise your hand and your designated AI ToT will assist you immediately.

03 · PATH B

Launch & Create First Project

Strike while the iron is hot. Initialize your workspace.

01

Launch App

Open Google Antigravity. Ensure your system meets the required permissions for standalone operation.

02

Add Folder

Click the **Add Folder** icon in the left sidebar to set up your isolated workspace project.

03

Start Agent

Select **Local/Worktree** mode in the chat input and boot up your first agent.



AI Tutors on Standby — Stuck on creating a project? Raise your hand and your designated AI ToT will assist you immediately.

What is GStack?

One install. **28 slash commands.**

Your coding agent becomes your entire team.

- **CEO**
challenges your product scope
- **Staff eng**
catches production bugs before you
- **Security officer**
runs OWASP + STRIDE
- **Designer**
builds design systems + variants
- **QA**
opens a real browser and tests your app
- **Release engineer**
ships, monitors, and deploys

Open-source skill pack by Garry Tan (Y Combinator). Built for Claude Code; installs into Codex with `./setup --host codex`.

WATCH

GStack Master Class — Garry Tan

How the Y Combinator CEO ships with an AI team. Watch, then we install.

WHILE YOU WATCH, LOOK FOR

- Which **role** the AI plays at each step (CEO, designer, QA...)
- Where he **stops the agent** and makes it think first
- How many **parallel sessions** he runs — and why

"Better tools demand better thinking." — the whole talk is that one sentence, demonstrated.

IF THE PLAYER IS BLOCKED

▶ **Watch on YouTube** youtu.be/wkv2ifxPpF8

Source: Garry Tan, GStack Master Class ·
youtube.com/watch?v=wkv2ifxPpF8

REALITY CHECK

Sharp planning is critical.

Better tools demand better thinking.

WHAT

What are we building?

WHY

Why is it needed?

HOW

How do we solve it?

WHO

Who is it for?

WHERE

Where do we start?

WHEN


When is it done?

GStack relentlessly challenges your thinking!

GSTACK

The Sprint Process

Each skill feeds the next. Nothing falls through the cracks.



/retro + /learn + /document-release → feeds back into next sprint



Why this matters for your startup: `/office-hours` asks the 6 forcing questions YC partners would ask. Challenges your assumptions before you write code.

04

Install GStack

Bring the tools to the forge. Clone the repository.

```
terminal - gstack  
  
$ git clone https://github.com/garrytan/gstack.git  
$ cd gstack  
$ ./setup # Claude Code (default)  
$ ./setup --host codex # Path A: install into Codex
```



AI Tutors on Standby — Need help navigating the terminal or encountering git errors? Raise your hand and your designated AI ToT will assist you.

05

Why Office Hours?

Raw ideas break under pressure. **The forge demands precision.**

THE RAW ORE · VAGUE IDEA

"Build an app for students."

→
/office-
hours

THE FORGED BLADE · REFINED

[Problem] Campus meal-sharing platform

[Audience] Korea University students

[Constraints] Dietary restriction filter, 15-min pickup, split group payment

[Metrics] Target: 50 completed orders/day

05

Bad vs Good Prompts

Real student examples. Precision dictates the output.

× THE VAGUE PROMPT

"Make a food delivery app"

Produces generic garbage.

✓ THE FORGED PROMPT

"Build a campus meal-sharing platform for Korea University students with dietary restrictions filter, 15-min pickup guarantee, and split-payment for group orders"

Produces precise solution.

05

The Refinement Process

AI demands clarity. Forge your raw ideas through these 5 precise stages.

01	PROBLEM	What exact pain are you solving?
02	AUDIENCE	Who specifically feels this pain every day?
03	CONSTRAINTS	What are the technical and operational limits?
04	METRICS	How will you empirically measure success?
05	VALIDATION	Why will people actually abandon their current habits to use this?

Mission 01: Office Hours

Validate your startup idea with an adversarial AI advisor.

COMMAND

/office-hours

EXPECTED OUTCOME

Refined pitch and a reframed problem ready for the forge.



Shape the idea before you swing the hammer.

Let the AI find the fractures in your plan before you write a single line of code.

05 · EXAMPLE 01

Example 01: The Interrogation

Raw ideas break under pressure. The AI advisor finds the fractures.

office-hours — session

```
$ /office-hours "I want to build a food sharing app for students"
```

```
[AI ADVISOR] IDEA TOO BROAD. The forge requires specifics.
```

1. Which students? (Commuters? Dorms? Dietary limits?)
2. What exact pain are you solving that existing apps ignore?

```
$ Target: Korea Univ students with dietary restrictions (vegan/halal).
```

```
$ Problem: Can't meet minimum delivery limits. Need group split-payments.
```

```
[REFINED PROMPT GENERATED]
```

```
"Build a campus meal-sharing platform for KU students with vegan/halal filters, 15-min pickup guarantee, and split-payment for group orders."
```

Tech Stack & MVP Scope

AI Advisor negotiates realistic constraints.

/office-hours SESSION

AI Advisor: What is your team's current tech stack? What is the timeline? Are there any hard constraints?

Student: We know some JavaScript. We have 2 weeks until demo day. Zero budget for servers.

AI Advisor: Understood. Scoping down... We must optimize for speed and JS familiarity. Forging technical constraints.

REFINED MVP SCOPE

Frontend: React Native (Expo)

Rationale: Leverages existing JS knowledge for rapid cross-platform deployment.

Backend: Firebase (BaaS)

Rationale: Zero upfront cost, requires no infrastructure management.


Timeline: 2-Week Sprint

Rationale: Cuts 'nice-to-have' features to ensure a working demo by the deadline.

HOW GARRY WORKS

The Conductor Pattern

10–15 parallel Claude Code sessions. Each in its own isolated workspace.

**GARRY'S NUMBERS · MARCH 2026**

- 3 projects simultaneously
- 15 concurrent conductor sessions
- 17,000 lines of code per day
- 35% of shipped code is tests
- 600K+ lines in 60 days (part-time)

HOW I ACTUALLY WORK

*"I don't have a to-do list anymore. I click + in Conductor. It creates a new worktree. Each one IS the to-do item. When I have an idea, I drop it in, run `/office-hours`, and start fleshing it out. It goes ding — my robot is ready. **I have 20 of these going all the time.**"*

05

Office Hours Output

The forged blueprint. No vague ideas, only precise execution constraints.

01 // Refined Problem Statement	Students with strict dietary restrictions (vegan, halal, allergies) spend avg. 45 mins coordinating safe group meals, leading to social exclusion.
02 // Target Audience	University undergrads (ages 18–24) who regularly eat lunch/dinner in groups of 3+ and have specific dietary constraints dictating vendor choice.
03 // Core Features (MVP scope)	1. Dietary Filter Engine (cross-references 5 major allergens/diets with campus vendor menus) · 2. Split-Payment Group Cart (one-click divided payments via KakaoPay API) · 3. 15-Minute Campus Pickup Guarantee (prioritized queueing with 3 anchor vendors)
04 // Success Metrics	50 daily active group orders within week 2. Less than 2% order cancellation rate due to dietary mismatch. High retention (3+ orders/week).
05 // Immediate Action	Execute <i>/design-shotgun</i> focusing purely on the group cart and filter UI. Ignore the admin dashboard for now.

/PLAN-CEO-REVIEW

Should we build this?

WHEN TO USE Before writing a single line of code.

REAL EXAMPLE · TELEMETRY SPRINT

19 bugs caught before a single line was written.

The agent tried to skip ahead to coding. Forced it to pause and meticulously run the CEO review. This caught critical misalignments in priority and prevented building the wrong features entirely.

WITHOUT REVIEW

- ✗ Wrong Problem
- ✗ Vanity Metric
- ✗ Wrong Priority

WITH /PLAN-CEO-REVIEW

- ✓ Validated Pain Point
- ✓ Actionable KPI
- ✓ Ruthless Focus

Tests verify you built what you planned. **Reviews verify you planned the right thing.**

REVIEW 02

Engineering Review > /plan-eng-review

WHEN TO USE After CEO review, to ensure technical architecture can be shipped safely.

WITHOUT REVIEW · THE RISKS

✗ Schema Conflicts

Executing an ADD and RENAME on the same column crashes the database.

✗ Migration Ordering

Creating a table before its foreign key dependency exists.

✗ Dependency Hell

Updating a core package that silently breaks downstream services.

WITH /plan-eng-review

✓ Conflict Caught

AI identifies the ADD/RENAME collision before a single line is written.

✓ Safe Sequencing

Proposes a safe, multi-step migration plan to preserve data integrity.

✓ Tree Validation

Dependency tree is verified and locked prior to implementation.

Catch infrastructure bugs before they crash production.

REVIEW 03

Codex Review

> codex review

The microscopic inspector. Finding the bugs you didn't know you wrote.

1 // Missing Call Sites

API endpoint updated on backend, but 3 frontend usages were forgotten and would crash.

2 // Edge Cases

Null checks, undefined states, and unhandled boundary conditions that humans routinely overlook.

3 // Off-by-One Errors

Flawed loop termination conditions and dangerous array index boundaries.

codex-inspector.log

⚠ **WARNING 1:** Unhandled Edge Case (Null Check)

- `const userName = user.profile.name;`

+ `const userName = user?.profile?.name ?? 'Guest';`

⚠ **WARNING 2:** Missing Call Site

API `'/v1/orders'` signature changed to require `'tenant_id'`.

✗ Found 3 frontend usages missing the new parameter.

✓ Auto-patching call sites in `Dashboard.tsx...`

✓ Auto-patching call sites in `Checkout.tsx...`

KEY TAKEAWAY: AI catches what humans miss in the details.

/PLAN-CEO-REVIEW

Business-Level Concern Detection

WHEN TO USE Before over-engineering a solution.

REAL EXAMPLE · ADMIN DASHBOARD SPRINT

Saved 2 weeks of engineering time by challenging assumptions.

The agent planned a complex React admin panel. The CEO review detected that the ops team only needed a simple weekly CSV export. We scrapped the UI and built a 1-hour cron job instead.

WITHOUT REVIEW

- ✗ Complex React UI
- ✗ 2 Weeks Effort
- ✗ Over-engineered

WITH /PLAN-CEO-REVIEW

- ✓ Automated CSV
- ✓ 1 Hour Effort
- ✓ Lean Solution

Code is a liability. The best PR is the one you don't have to write.

REVIEW 02

Engineering Review > /plan-eng-review

WHEN TO USE Validate API contracts and ensure backward compatibility before shipping.

WITHOUT REVIEW · THE RISKS

✗ Breaking Contracts

Changing a payload field from String to Object crashes existing mobile clients instantly.

✗ Missing Versioning

Overwriting an active endpoint instead of creating a v2 forces uncoordinated client updates.

✗ Deprecation Blackout

Dropping legacy auth tokens without a sunset period locks out older application versions.

WITH /plan-eng-review

✓ Contract Validation

AI detects payload mutations and flags potential breakage for downstream consumers.

✓ Semantic Routing

Proposes a strict v2 routing strategy while maintaining a backward-compatible v1 adapter.

✓ Sunset Strategies

Drafts a safe deprecation timeline and header-based warning system for legacy clients.

Protect existing clients. Never break backward compatibility silently.

REALITY CHECK


GARBAGE IN, GARBAGE OUT.

AI doesn't build apps. It builds exactly what you tell it to build.

A vague business idea is a blunt tool. You must forge your thoughts before generating code.


Start with **Why.**

Simon Sinek · TEDxPugetSound, 2009



IF THE PLAYER IS BLOCKED

▶ **Watch on YouTube** youtu.be/qp0HIF3SfI4



Buy the book · Amazon amazon.com/dp/1591846447 →

Design Shotgun


Cast multiple functional prototypes in parallel. Strike while the iron is hot.

COMMAND

/design-shotgun

EXPECTED YIELD

**3 to 5 distinct UI mockups forged instantly,
ready for visual selection.**



Code Generation

Hand over the blueprints. AI builds the app layer by layer.

PROMPT · IN YOUR AGENT

```
> Implement the approved plan.  
> Start with the group-cart screen. Stop  
and show me before each new file.
```

No special slash command here — after `/autoplan`, you simply tell Codex or Antigravity to build. Keep the plan file open.

EXPECTED OUTCOME

Functional codebase ready for testing.

The forge never sleeps. Observe the terminal output.

MISSION 04

QA Automation

Test everything with **browser automation**.

COMMAND

/qa

EXPECTED OUTCOME

**Automated test suite running via Playwright.
Zero regressions before shipping.**

playwright / test-runner

Test Execution Pipeline

✓ auth.spec.ts — 1.2s

✓ checkout.spec.ts — 3.4s

● api_integration.spec.ts — **running...**

EVERYTHING STARTS HERE

Open the **LMS** now.

All lecture material, labs and daily submissions live on this one page. Scan it, or type it — it is short on purpose.

TYPE THIS IN YOUR BROWSER

da.gd/**6H8Fd**

Case matters: capital **H**, capital **F**, the rest lowercase.



KEEP FORGING AFTER TODAY

GitHub

github.com/garrytan/gstack

Documentation

antigravity.google/docs

Community

Discussions & forum

Master Class

youtu.be/wkv2ifxPpF8